


NanoLog: A Nanosecond Scale Logging System

武新楠

M201877271



HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Overview

- **Implemented a fast C++ Logging System**
 - NanoLog是一种新的日志系统,比现有的解决方案快1-2个数量级
 - 像Log4j2一样,保持了类似printf的API
 - 吞吐量高达每秒8000万个日志消息
- **Shifts work out of the runtime hot-path**
 - 将大部分工作从程序运行时转移到程序编译和执行后阶段
 - 通过编译时提取静态内容来减少用户日志调用
 - 重写日志代码仅在运行时发出动态信息
 - 依赖后处理器来重新整理日志

What is Logging?

- **printf-like messages with dynamic information**
 - Written at development time; output at runtime
 - Developer specifies severity, log message + dynamic data
 - Logging system inserts invocation date/time, file location

```
main.cc
1 void main() {
2     LOG(NOTICE, "The task began");
3     float result = doSomething();
4
5     LOG(NOTICE,
6         "The result was %4.2f", result);
```

The diagram illustrates the structure of a log message. It shows two example log entries: "2017-11-17 21:35:16 main.cc:3 NOTICE[4]: The task began" and "2017-11-17 21:38:29 main.cc:5 NOTICE[4]: The result was 12.00". Red brackets are used to identify the components of these messages. The first three components (Date/Time, File Location, and Severity) are grouped under a bracket labeled "Context (Implicit)". The final component (the log message) is grouped under a bracket labeled "Message (explicit)".

Date/Time	File Location	Severity	Message
2017-11-17 21:35:16	main.cc:3	NOTICE[4]	The task began
2017-11-17 21:38:29	main.cc:5	NOTICE[4]	The result was 12.00

What makes logging slow?

```
1473057128.133777014 src/LogCleaner.cc:826 in TombstoneRatioBalancer  
NOTICE: Using tombstone ratio balancer with ratio = 0.400000
```

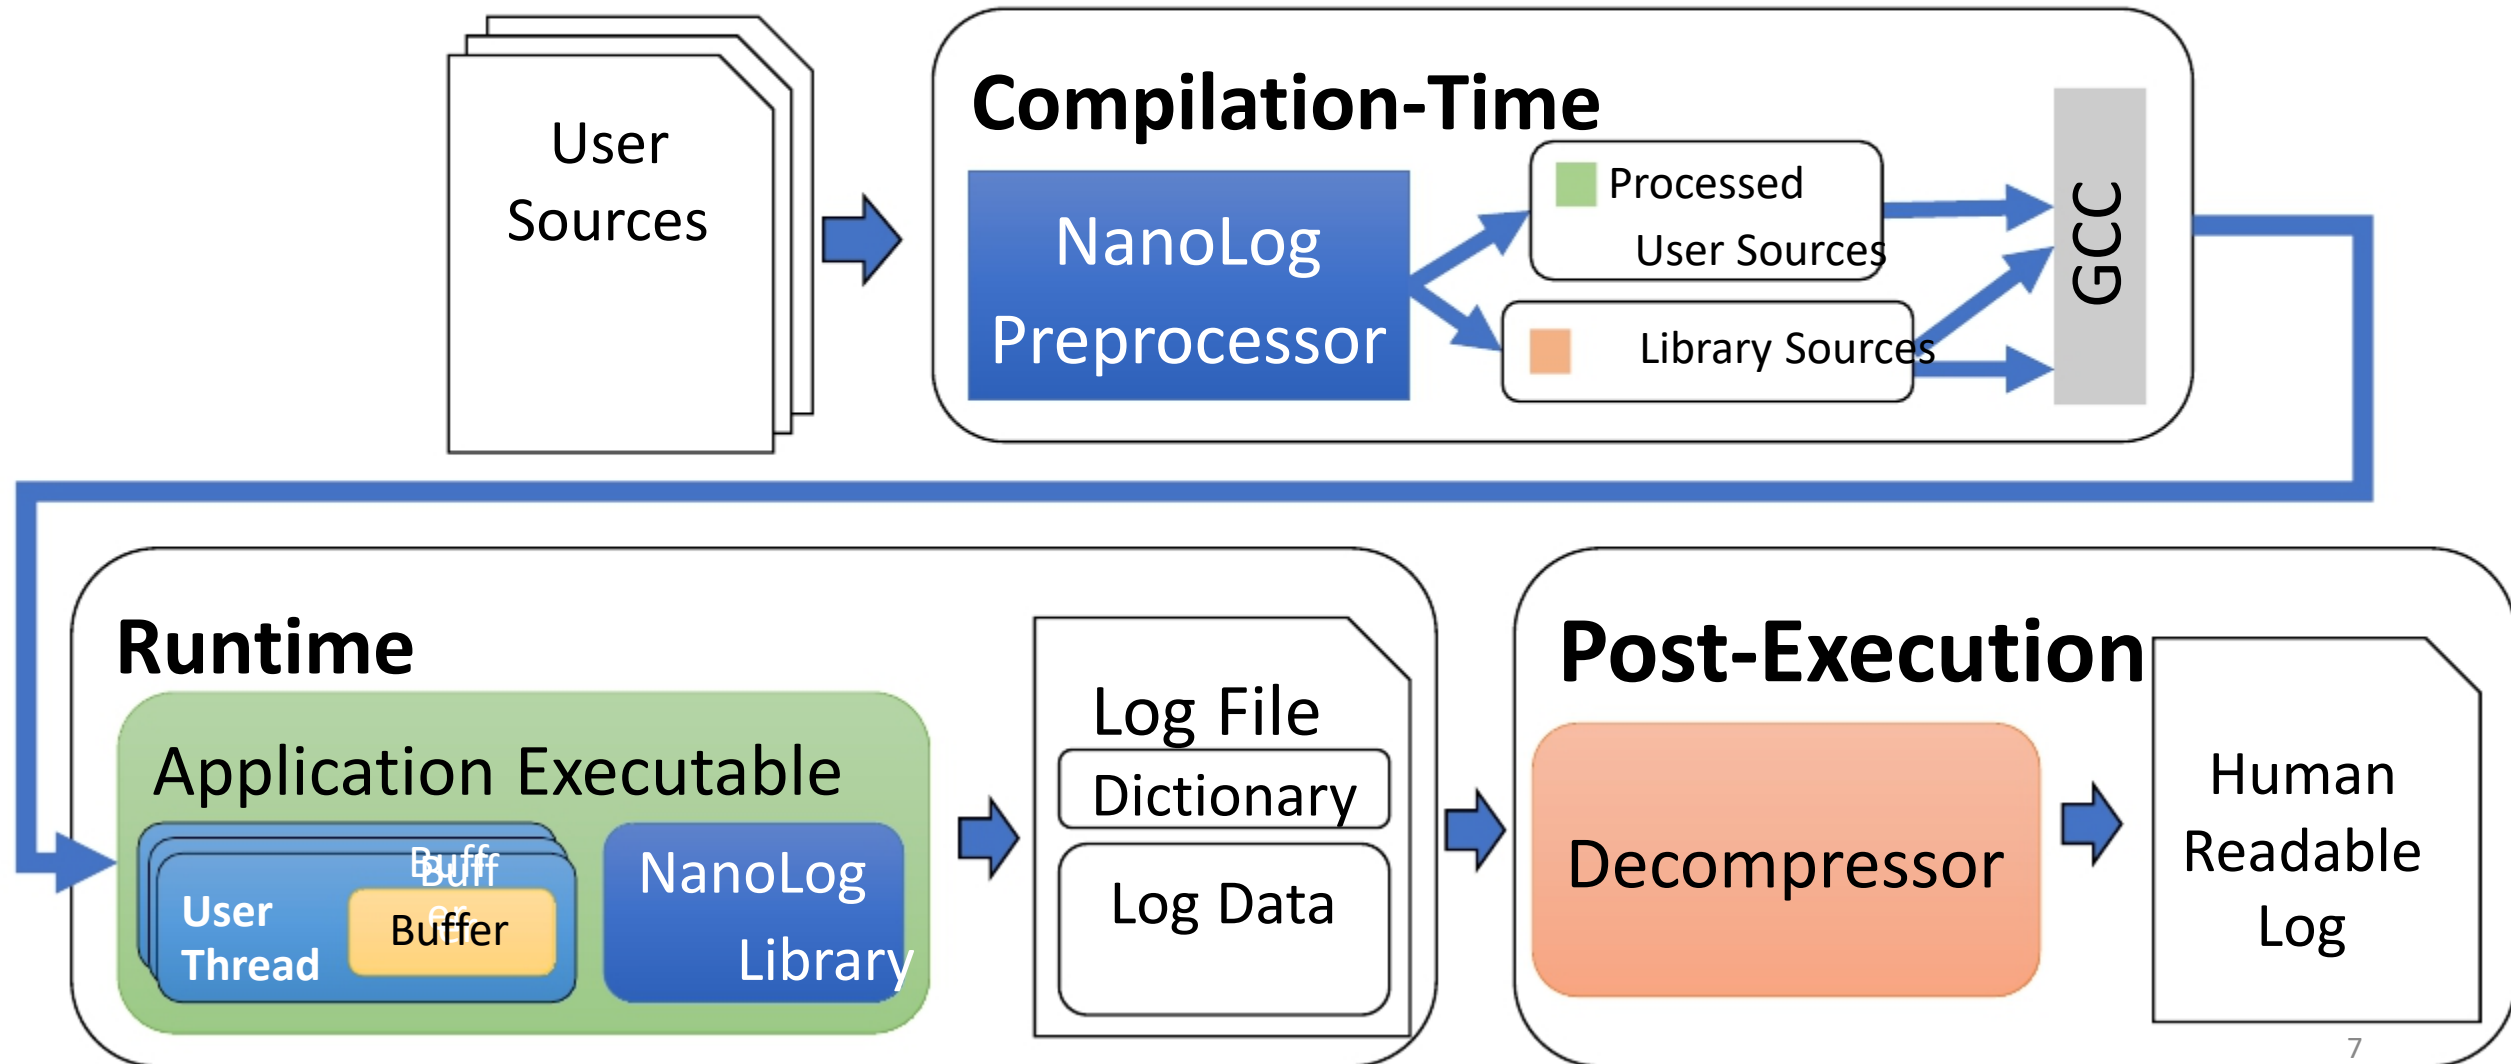
- **Compute: Complex Formatting**
 - Loggers need to provide context (i.e. file location, time, severity, etc)
 - The message above has **7 arguments** and takes **850ns** to compute
- **I/O Bandwidth: Disk IO**
 - On a 250MB/s disk, the **129 byte** message above takes **500ns** to output!

Solutions

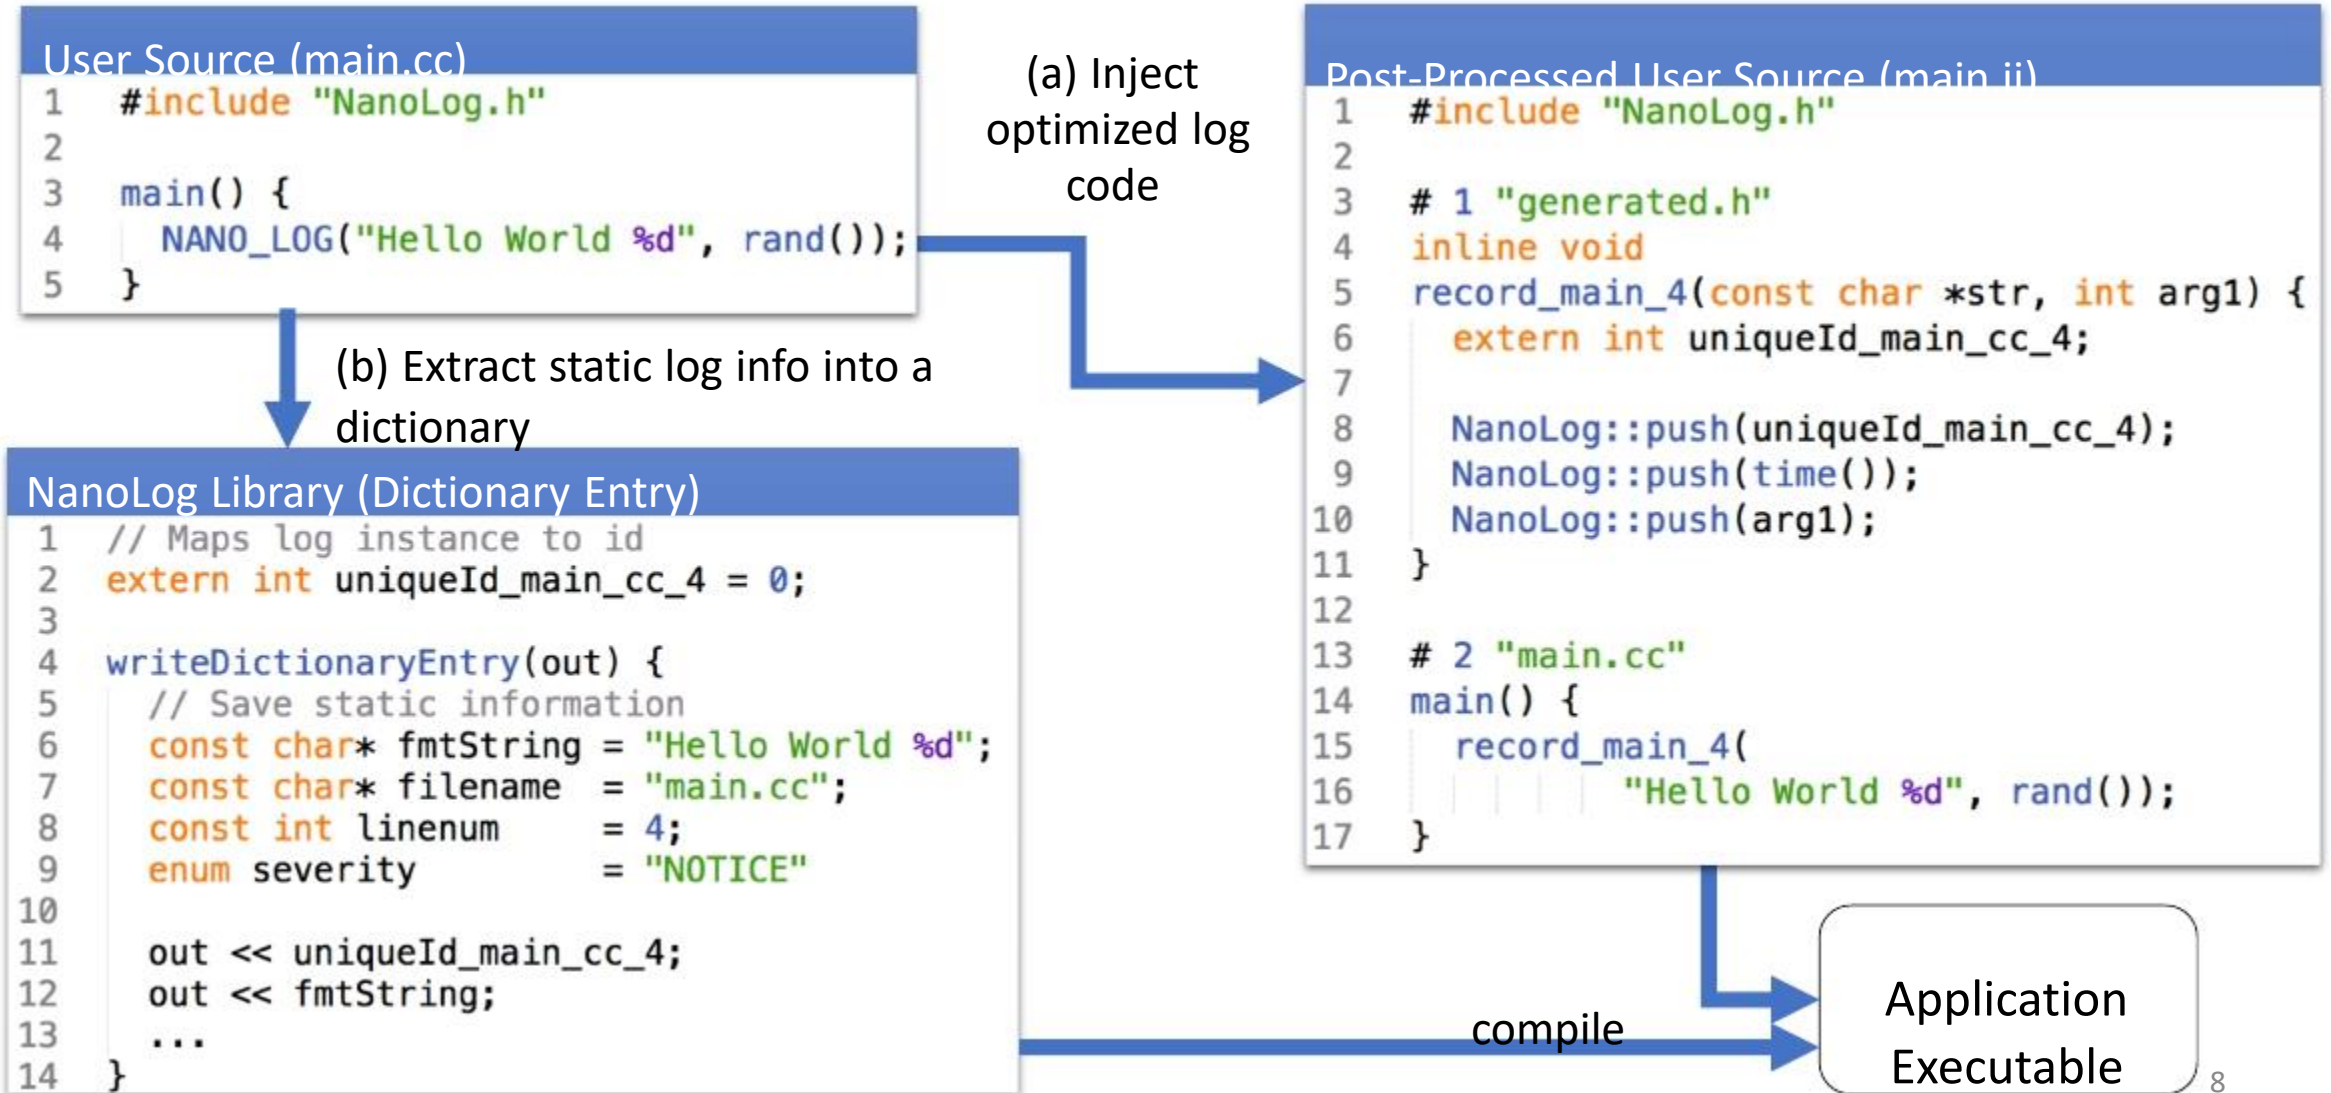
```
1473057128.133777014 src/LogCleaner.cc:826 in TombstoneRatioBalancer  
NOTICE: Using tombstone ratio balancer with ratio = 0.400000
```

- **Compute: Defer formatting to an offline process**
 - Developers only look at a small portion of the log
 - Seems wasteful to output a human-readable log eagerly
- **I/O Bandwidth: De-duplicate Static Information**
 - Log messages contain a lot of information known at compile-time
 - i.e. file location, line #, function, severity, format string
 - Logging only the dynamic information in binary saves I/O
 - Shrinks the 129 byte message above to just 16 bytes

NanoLog System Architecture

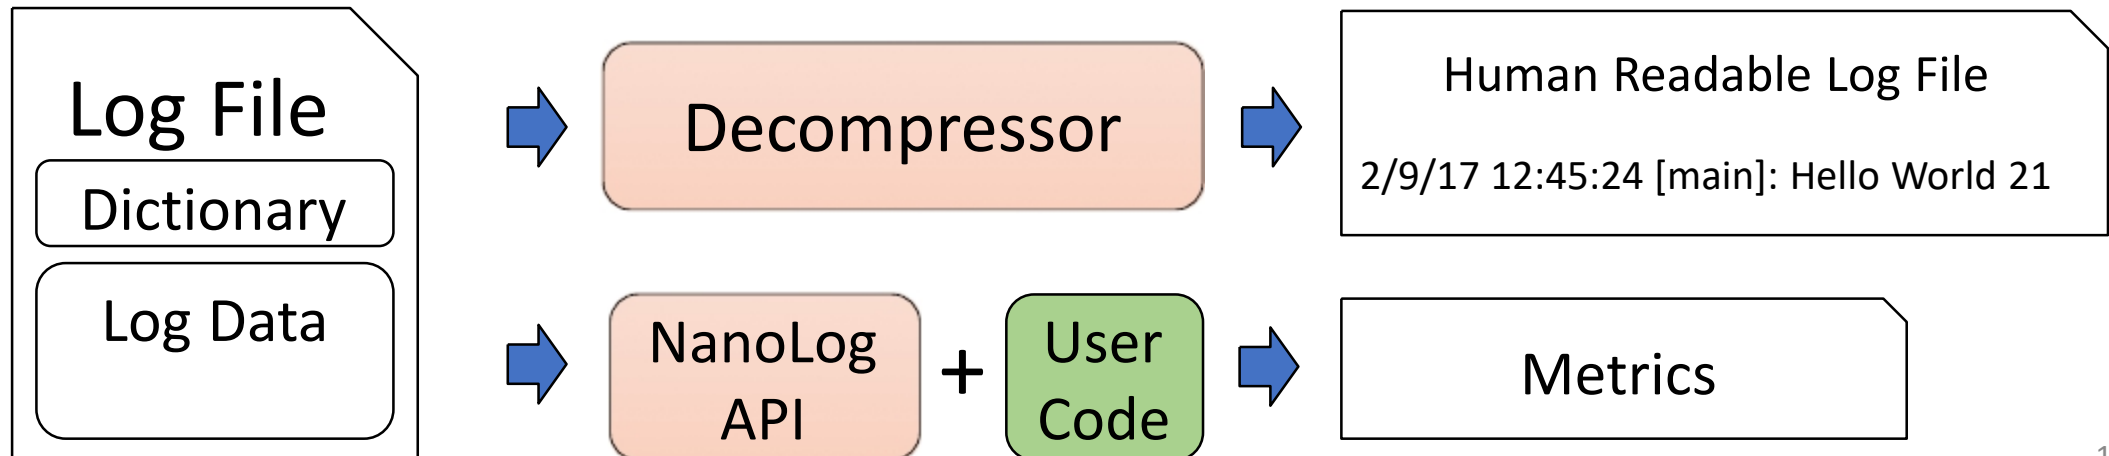


Compile-time Optimizations(编译过程)




Decompressor(解压器)

- **Reconstitutes the log data by combining static + dynamic info**
- **Stand-Alone**
 - Reads the log file and produces a full human-readable log file.
 - Ultimately pays the formatting + output costs of the full log
- **Programmatic API (for fast aggregation)**
 - Process the log messages one by one without formatting.
 - More efficient than processing data in ASCII



Microbenchmark - Throughput



Evaluation

- Repeated logged 1 message with no inter-log delay in a thread
- Varied the number of logging threads to maximize throughput

Conclusions

- NanoLog always faster
- Even better with fewer dynamic arguments.

ID	ExampleOutput	NLMsgSize
staticString	Startingbackupreplicagarbagecollectorthread	3-4Bytes
stringConcat	Openedsessionwithcoordinatoratbasic+udp:host=192.168.1.140,port=12246	43Bytes
twoIntegers	bufferhasconsumed1032024bytesofextrastorage,currentallocation:1016544bytes	10Bytes


underlined indicates dynamic data¹²

Microbenchmark - Unloaded Latency

- **Goal: Measure minimum processing delay**
- **Setup:**
 - 100,000 iterations of logging 1 log message with a 650ns delay between each to eliminate blocking due to I/O
 - Percentile times reported below in nanoseconds

System	NanoLog	spdlog	Log4j2	ETW		NanoLog	spdlog	Log4js	ETW
LogMsgs	MedianLatency(ns)					99.9PercentileLatency(ns)			
staticString	8	230	192	180		33	473	1868	726
stringConcat	8	436	230	208		33	1614	6171	2954
twoIntegers	7	674	160	200		44	1335	1992	761

NanoLog For Analytics



Mock Analytics Workload

1. Find all instances of “Hello World # %d”
2. Interpret “%d” as an integer and compute the min/mean/max

Configurations

- **NanoLog:** Used the compact, binary log
- **C++:** Used *atoi()* on the full log
- **Python, Awk:** Use regex on the full log

NanoLog Analytics Benefits:

- Smaller log files (~747MB vs. 7.6GB)
- Binary format (no ASCII parsing req.)

Conclusion

- **NanoLog's Techniques**

总之,NanoLog是一种新的日志系统,比现有的日志系统快1-2个数量级。它通过使用预处理器在编译时删除静态日志信息,注入优化代码,仅记录运行时的动态信息,并将日志消息的格式化推迟到程序执行后来加速日志记录。

- **Benefits:**

该系统的好处在于他运行时非常高效,并且可以在执行后提供快速聚合。唯一需要注意的是,它可能不会立即以人类可读的格式显示日志。

Thanks!